

Literature Review in Computational Game Theory Topics

Games of Unawareness

- **Interactive Unawareness:** *Heifetz, Meier & Schipper (2006)* introduced a formal state-space model allowing players to be unaware of some game elements ¹. This framework was refined in a canonical model by the same authors ², enabling analysis of games where different players might **not conceive** of certain moves or states. Subsequent work extended these ideas to dynamic games and speculative trade under unawareness ³.
- **Knowledge and Unawareness:** *Halpern & Rêgo (2008, 2009)* developed logical approaches for reasoning about knowledge of one's own unawareness ⁴ ⁵. They later modeled extensive-form games with possibly unaware players, defining equilibrium concepts when players may discover new actions over time ⁶. *Li (2009)* similarly studied information structures with unawareness in Bayesian games ⁷.
- **Equilibrium Under Unawareness:** Given the diversity of frameworks, equilibrium notions vary. For example, *Feinberg (2012)* unified normal, extensive, and repeated games of unawareness, defining Nash-like solutions and showing they yield novel behaviors beyond standard incomplete-information games ⁸ ⁹. *Schipper & Meier (2014)* introduced “**unawareness perfection**” for Bayesian games, refining equilibrium by requiring stability under zero-probability awareness events ¹⁰.
- **Surveys and Applications:** *Schipper (2014)* provides a brief survey of the unawareness literature ¹¹, comparing epistemic models from economics and computer science. The literature has grown to include contract theory (e.g. **incentives for unaware agents** ¹²), trade with unaware consumers ¹³, and behavioral experiments on how exposure to unawareness affects decisions ¹⁴. This special issue survey highlights that unawareness challenges classical game assumptions, requiring new solution concepts and logical tools ¹⁵.

Sparsification of Equilibria (Sparse/Low-Rank Strategies & Norm-Regularization)

- **Equilibria with Small Support:** It is known that even though Nash equilibria exist, they may have full (“dense”) mixed strategy support, complicating interpretation ¹⁶. However, if a two-player game's payoff matrices have *low rank*, then an exact Nash equilibrium exists with **small support** (few actions per player) ¹⁷. This result (from algorithmic game theory) implies that structural payoff constraints yield simpler equilibria that are easier to compute and communicate. In general bimatrix games, any ϵ -equilibrium can be approximated by a strategy mixing over relatively few actions ¹⁸ ¹⁹, though finding the minimal support is NP-hard.
- **k -Sparse Strategy Commitment:** Recent work explicitly studies equilibria where one player commits to a mixed strategy using at most k actions. *Afiouni et al. (2025)* examine **k -sparse commitments**, proving structural properties and proposing algorithms for Stackelberg and zero-sum games ¹⁶ ²⁰. Their methods (mixed-integer programming plus iterative strategy generation) attain over 90% of the full Nash payoff with very small supports, demonstrating that sparse approximate equilibria often suffice in security games and other applications ²¹. These

results underscore the trade-off between optimality and simplicity: even modest support limits can drastically improve human interpretability and reduce coordination complexity.

- **Norm-Regularized Equilibria:** Drawing inspiration from compressive sensing, some studies seek **well-1\$ or other norm-regularization** to induce sparse equilibrium strategies ²² ²³ . While one cannot directly regularize the Nash equilibrium conditions (they are combinatorial ²³), regularization ideas have influenced heuristic methods to select an equilibrium with minimal support or to limit coordination among team members. For example, in team/cooperative games, one might impose communication or complexity costs that effectively penalize highly mixed strategies. *Lipton, Markakis & Mehta (2003)* famously showed any 2-player game has an ϵ -Nash equilibrium with support $O(\log n)$ (for n actions) by randomization arguments, hinting that **equilibrium sparsification** is generally possible with small welfare loss.
- **Low-Coordination Team Strategies:** In team games (common-payoff scenarios), researchers have explored how to achieve near-optimal outcomes with limited ex-ante coordination. One approach restricts the strategy space to **low-rank policies** — essentially limiting the dependence between teammates' actions. This connects to work on correlated equilibrium with restricted signals (e.g. *low-dimension correlation devices*). Although specific literature is sparse, the intuition is that if team members agree on a simple joint plan (low communication), the equilibrium might sacrifice little efficiency. For instance, *Björnerstedt & Weibull (1996)* showed that even boundedly rational (simple) strategies can perform well in team coordination games. Modern multi-agent learning uses population-based training to maintain a **diversity of styles at high skill** ²⁴ , effectively ensuring agents learn robust strategies requiring minimal post-hoc coordination. This remains an open research direction: balancing performance and simplicity (or norm-based “regularity”) in multi-agent equilibrium selection.

Games with Stochastic Action Sets

- **Stochastic Action Availability:** Traditional game models assume fixed action sets, but in many settings the *available actions themselves are random*. For example, in online platforms the choices (ads, recommendations) vary stochastically on each encounter ²⁵ . *Boutilier et al. (2018)* formalized this as a **stochastic action set MDP (SAS-MDP)** ²⁶ . They showed that when the set of feasible actions at a state is a random variable, one can still compute optimal policies by augmenting state with action-set information. In particular, they proved that the Q-value of an action is independent of which other actions are available ²⁷ , enabling efficient learning and planning despite a combinatorially varying action set. This foundational work provides algorithms for SAS-MDPs (both model-based and Q-learning variants) and demonstrates significant benefits in applications like recommender systems and ad placement ²⁵ .
- **Sleeper Arms and Contextual Bandits:** In online learning, this scenario appears as **sleeping bandits** or **contextual bandits with random arm availability**. Algorithms have been developed for bandits where at each round a random subset of arms is available (e.g. due to inventory or scheduling) ²⁸ . These works (Kleinberg et al. 2010; Kanade et al. 2014) provide regret bounds assuming various stochastic availability patterns ²⁹ ³⁰ . Overall, they show that efficient learning is still possible if the agent properly accounts for the probability of each action being available, often by treating unavailability as “missing data” rather than failure.
- **Game-Theoretic Modeling via Nature:** In game theory, a conceptually equivalent approach is to add a **chance player (“moves by nature”)** that randomly selects which actions are available for the strategic players ³¹ . This transforms a game with stochastic action sets into a larger extensive-form game with imperfect information. As noted by Osborne & Rubinstein (1994), any stochastic outcome (including random menu of actions) can be modeled by an additional player that moves first with a given probability distribution ³² ³¹ . Using nature's moves, one can define equilibria in these games: players choose optimal strategies given the random constraints

on their actions. This method has been applied, for instance, in **stochastic games** (Shapley 1953) and **Bayesian games** where the type space can encode random action availability. While the general equilibrium analysis becomes complex (players form beliefs about which actions might be available), the approach provides a sound theoretical foundation – essentially reducing stochastic action sets to a special case of Bayesian uncertainty.

- **Applications:** Games with stochastic action sets appear in network security (where certain defense actions may or may not be feasible at a given time), finance (investment opportunities come and go randomly), and multi-robot planning (robots may face random failures of certain moves). The key challenge is ensuring strategies are robust against such random constraints. Recent research in multi-agent reinforcement learning also considers **dynamically changing action spaces** (e.g. new agents or actions entering a game), and leverages techniques like meta-learning to adapt policies when the available action set shifts over time. This is an emerging area at the intersection of game theory and AI, with many open problems about equilibrium existence and computation when each stage's action set is drawn from a distribution.

Payoff Set Learning and Characterization

- **Inverse Game Theory:** A fundamental question is how to **learn the payoff functions or payoff sets** of a game from observations of strategic behavior. *Vorobeychik, Wellman & Singh (2005)* pioneered this by treating payoff learning as a regression problem ³³ ³⁴. Given sampled strategy profiles and observed payoffs, they fit parameterized payoff functions that explain the data, focusing on games with infinite strategy spaces. Their approach leverages structure (e.g. symmetry) to generalize from finite data, and evaluates success by how well the learned game predicts agents' equilibrium strategies ³⁵ ³⁶. This line of work essentially performs **inverse reinforcement learning** in a multi-agent setting: deducing what underlying payoffs would make the observed behavior rational.
- **Feasible Payoff Set Inference:** More recently, researchers consider scenarios where we observe an equilibrium outcome (or several) and wish to infer all games consistent with those outcomes. *Fabricius et al. (2026)* formulate the problem of recovering the **feasible payoff set** given that players are (approximately) in equilibrium ³⁷. For example, if two players are playing a Nash equilibrium of some unknown game, what constraints does this impose on the payoff matrices? By observing action frequencies or mixed strategies, one can derive a set of payoff matrices for which the observed strategy profile is an ϵ -Nash equilibrium ³⁸ ³⁹. The goal is to learn this set (or a tight approximation) with high confidence. Initial results provide sample complexity bounds for estimating payoff sets under different equilibrium assumptions (exact vs. approximate, zero-sum vs. general-sum) ⁴⁰ ⁴¹. This research merges statistical learning with game-theoretic solution concepts, aiming to **identify games from behavioral data**.
- **Characterizing Equilibrium Payoff Regions:** In repeated and stochastic games, a classic result is the **Folk Theorem**, which characterizes the set of achievable equilibrium payoffs. For infinitely repeated games with discounting, any payoff profile that exceeds the players' minmax (security) levels and is feasible can arise as a Nash equilibrium payoff (Fudenberg & Maskin, 1986). Computationally, recent work has tried to *learn* or approximate these equilibrium payoff sets. For instance, algorithms exist to compute the set of Pareto-efficient equilibrium payoffs in two-player repeated games, by iteratively learning best responses and adjusting punishments. Additionally, in mechanism design and multi-agent learning, researchers examine how agents can learn to cooperate or collude to realize certain payoffs (which amounts to exploring the equilibrium payoff frontier).
- **Learning in Games:** On the algorithmic side, *Brafman & Tennenholtz (2003, 2004)* explored **learning equilibria**: conditions under which agents' adaptive algorithms will converge to an equilibrium of the game. In coordination games, they gave model-based algorithms that efficiently identify equilibrium strategies via exploration ⁴². In general-sum games, they

introduced the notion of an **efficient learning equilibrium** – a state where each agent’s learning process stabilizes to an equilibrium strategy profile, and no one can gain by deviating from their learning rule ⁴³. These works imply that, under certain restrictions, multi-agent systems can converge to equilibrium behavior while simultaneously *estimating* the game’s payoff structure. Such results connect to the idea of “self-play” in AI: agents can learn the game and the optimal strategies at the same time.

- **Applications:** Payoff learning and characterization is critical in economic domains where the payoff structure is unknown or changing – for example, inferring competitors’ utility functions in auctions or estimating drivers’ payoff matrices in traffic routing games. By learning payoff sets, one can predict how behavior might change if the game’s parameters shift. It also aids mechanism designers in verifying if observed outcomes correspond to equilibrium of some intended game (for instance, validating that an auction produced nearly efficient outcomes, implying bidders’ values aligned with expectations). As AI agents become more prevalent, **interpreting observed strategic interactions** in terms of underlying incentives will be an important tool for transparency and regulation.

Polymatroid Regret Minimization, Online Learning, and Gittins Index

- **Multi-Armed Bandits and Indices:** The **Gittins index** is a fundamental concept for decision-making under uncertainty, especially in multi-armed bandit problems. *Gittins (1979)* showed that for a single player facing multiple independent reward processes (arms), there exists an index for each arm such that always pulling the arm with the highest index is optimal ⁴⁴. This index policy maximizes expected reward in a wide range of settings and converts a multi-armed problem into a series of one-arm decisions. *Whittle (1980)* further elucidated the properties of the Gittins index ⁴⁴, and it has since been applied to queueing systems, scheduling, and search problems. A recent tutorial by Scully & Terenin (2025) demonstrates the practical use of Gittins indices in complex domains (like Bayesian hyperparameter optimization), underscoring that while its theoretical definition is subtle, it can significantly outperform naive exploration strategies ⁴⁵ ⁴⁶.
- **Polymatroid Optimization and Indexable Systems:** Many stochastic scheduling problems can be cast as **extended polymatroids**, where the set of achievable performance vectors forms a polymatroid (a polyhedron with matroid structure). *Bertsimas & Niño-Mora (1996)* developed a polyhedral approach linking **conservation laws** in queueing to index-based policies ⁴⁷ ⁴⁸. They proved that for certain multi-class scheduling problems, the achievable region is a polymatroid and the vertices correspond to priority policies; thus an optimal policy is an extreme point found via a greedy algorithm (index policy) ⁴⁷ ⁴⁹. In particular, they unified many classical results by showing that if a system is *indexable* (each component’s propensity to be served can be indexed independently), then a greedy index rule (like Gittins or Whittle index) is optimal ⁵⁰. This framework not only explained why Gittins indices work for multi-armed bandits, but extended to restless bandits (via Whittle’s partial conservation laws) and other resource allocation problems.
- **Combinatorial and Structured Bandits:** In online learning, **combinatorial multi-armed bandits** generalize the bandit problem to structured action spaces (choosing sets of arms, paths in a graph, etc.). *Chen, Wang & Yuan (2013)* introduced a general CMAB framework with linear and nonlinear reward functions ⁵¹. They assume an offline optimization oracle for the combinatorial problem and derive UCB-type algorithms with regret bounds that scale with the structure’s complexity. For example, selecting k out of n arms (a size- k subset) each round can be done with $O(\log n)$ regret per round, given appropriate smoothness conditions ⁵² ⁵³. When the feasible sets form a matroid, specialized algorithms exploit greedy-optimality. *Kveton et al. (2014)* developed **Matroid Bandits**, proving that if the reward is a modular function on a

matroid, an optimistic greedy algorithm achieves sublinear regret and is computationally efficient ⁵⁴ ⁵⁵ . Their **Optimistic Matroid Maximization (OMM)** algorithm essentially runs a UCB on each element and uses a greedy selector to pick the best independent set, attaining near-optimal learning performance with time per round only $O(L \log L)$ for L actions ⁵⁵ ⁵⁶ .

- **Regret Minimization in Polymatroid Environments:** The notion of *polymatroid regret* appears when multiple resources or arms are allocated simultaneously under a global constraint. In such cases, one can often derive an **upper confidence bound algorithm** that respects the polymatroid structure, achieving regret that grows only sub-linearly in time. An example is dynamic resource allocation in wireless networks, where each timeslot's rate region is a polymatroid and the goal is to meet long-term throughput targets. Algorithms that treat this as an online linear optimization over a polymatroid (using mirror descent or exponential weights) have been shown to converge to the optimal operating point with low regret. These methods leverage the polymatroid's greedy optimization property (every linear objective is optimized at a vertex), which aligns with picking an extreme point each round akin to an index policy.
- **Connections to Coherence and Blackwell Approachability:** The term "coherence" in this context may refer to **Blackwell's approachability theorem**, which deals with ensuring a vector of long-run average outcomes stays in a desirable set (often a convex set like a polymatroid). Coherence implies that if a set of payoff vectors is approachable, a strategy exists to guarantee the time-average payoff converges to that set. This has been applied in regret minimization: for example, designing no-regret algorithms for vector-valued rewards by projecting the cumulative reward vector onto a cone or polytope (a form of ensuring no external regret across multiple objectives). In bandit problems, an algorithm is coherent with respect to a polymatroid if it treats each resource allocation consistently so as to not violate the underlying conservation laws.
- **Key References:** Classic references include *Gittins (1989)* for a comprehensive treatment of index theory, *Whittle (1988)* for restless bandits and index heuristic, and *Weber & Weiss (1990)* on exploring the empirical properties of index policies. For online learning, *Audibert, Bubeck (2010)* on multi-armed bandits, *Agrawal & Devanur (2014)* on bandits with concave rewards, and *Lattimore & Szepesvári (2020)* provide modern theoretical foundations. These works collectively bridge **game-theoretic planning and online learning**, showing that structured problems (those with polymatroid or matroid structure) allow efficient algorithms that carry strong performance guarantees (optimal indexes or low regret). The Gittins index remains a guiding principle, exemplifying how understanding problem structure (independence and decomposability) leads to powerful solutions in both stochastic control and online learning contexts ⁵⁰ ⁵⁷ .

LTL_f and Formal Methods in Games

- **Linear Temporal Logic on Finite Traces (LTL_f):** Standard LTL is interpreted over infinite plays (traces), but many games and planning problems are naturally finite-horizon. *De Giacomo & Vardi (2013, 2015)* introduced **LTL_f**, a variant of linear temporal logic for finite traces, and showed it can be used for specifying goals in single-agent and multi-agent settings ⁵⁸ ⁵⁹ . LTL_f has the same syntax as LTL (with operators X, F, G, U) but evaluated on finite sequences ⁵⁸ . It has expressiveness equivalent to first-order logic over finite words and is amenable to automata-based analysis similar to LTL ⁶⁰ ⁶¹ . Key distinctions include the *weak Next* operator (X_w) to handle end-of-trace semantics ⁶² . LTL_f has found applications in **automated planning and verification**, where one needs to reason about terminating scenarios (e.g. plan execution). Toolchains now exist for LTL_f satisfiability, model checking, and reactive synthesis ⁶³ ⁶⁴ .
- **Game Synthesis and DFA Games:** Synthesis from temporal logic specifications is essentially a two-player game problem: the system vs. the environment. For LTL_f specifications, the prevailing approach is to reduce to a **game on a DFA** (deterministic finite automaton). An

LTL $_f$ formula φ can be compiled into a DFA that accepts exactly the finite traces satisfying φ ⁶⁵. Then, a two-player game is constructed between System and Environment on the product of this DFA with the game arena ⁶⁶ ⁶⁷. A win for System corresponds to realizing the specification. Because the DFA is over finite words (rather than a Büchi automaton for infinite words), the game objective often becomes a simpler condition like reachability or safety in the product ⁶⁸ ⁶⁹. This means LTL $_f$ synthesis can avoid the complexity of parity games: for instance, with simple fairness assumptions, the winning condition can be solved via a nested fixed-point of depth 2 (a reachability within a safety) ⁷⁰ ⁷¹, as opposed to potentially many nested fixpoints for general LTL. Empirical studies show LTL $_f$ synthesis (using tools like **Syft** or **Nike**) often outperforms standard LTL synthesis in practice ⁷² ⁷³, because determinization of DFA over finite words is more tractable. Recent advances include on-the-fly LTL $_f$ synthesis that builds the DFA incrementally during the game solving ⁷⁴.

- **ATL and Strategy Logics:** In multi-agent systems, **Alternating-time Temporal Logic (ATL)** allows reasoning about strategic abilities: what coalitions can ensure a temporal property. *Alur, Henzinger & Kupferman (2002)* defined ATL (on infinite games), and complexity results show ATL model checking is PTIME and ATL * (with nested temporal operators) is 2EXPTIME-complete ⁷⁵. Recent work has considered ATL on finite horizons (ATL $_f$). Mogavero et al. (2014) and Murano et al. studied ATL $_f$ where path formulas are evaluated on finite traces. Model-checking ATL $_f$ can often be done via reduction to LTL $_f$ games: for instance, to check if a coalition A can enforce ψ until the end of the play, one constructs a game that combines the CGS (concurrent game structure) with the DFA for ψ ⁷⁶ ⁶⁷. This yields results like ATL $_f$ with perfect recall is PTIME for ATL $_f$ (analogous to ATL) and 2EXPTIME for ATL *_f (analogous to ATL *) ⁷⁷ ⁷⁸. These logics empower the specification of rich objectives in games – e.g. “agents 1 and 2 can cooperate to achieve Φ before time T regardless of what agent 3 does.” Strategy synthesis for ATL often involves automata or fixpoint computations similar to those for LTL but must account for the choices of multiple agents.
- **Model Checking and Verification:** Temporal logics (LTL, CTL, ATL) are widely used to **verify game properties**. For instance, one can ask if a given game state satisfies that “there is a strategy for Player 1 to eventually reach a goal state while avoiding hazards,” which is an ATL formula. Model checking that on a game graph amounts to solving a game with a temporal logic win condition. Tools from formal methods (like NuSMV, PRISM) can model check temporal specifications on game models (often called **turn-based** or **concurrent game structures**). The results mentioned in Table 1 and 2 of the ATL $_f$ paper ⁷⁵ ⁷⁷ summarize the complexity: CTL and ATL remain PTIME over finite traces, whereas allowing nested path quantifiers (CTL * , ATL *) raises complexity to PSPACE or EXPTIME, and adding imperfect information or memory can escalate it further (often undecidable in general). Nonetheless, for many practical cases (e.g., GR(1) fragment of LTL, or simple reachability objectives in ATL), **efficient model-checking algorithms** exist and are used in synthesis of controllers, protocol verification, and robot motion planning.
- **Game-Theoretic Analysis with LTL objectives:** There is also a growing literature on **reinforcement learning with temporal logic rewards**. Here an agent’s goal is given by an LTL/ LTL $_f$ formula instead of a scalar reward. The typical approach is to convert the formula into an automaton and augment the state space of an MDP or game with the automaton state (product MDP). Then one can apply RL or planning algorithms to maximize the probability of satisfying the formula ⁷⁹. For multi-agent competitive scenarios, this becomes a stochastic game with ω -regular (or finite-horizon temporal) objectives – essentially a **graph game with an LTL winning condition**. These are known to be decidable via automata-theoretic methods (Safra’s construction, etc.), but are 2EXPTIME-complete in general. Practical heuristics and GR(1)-synthesis approaches manage important subclasses efficiently.

- **Interpretability and Connections to Games:** Using formal specifications in games also improves interpretability of strategies. A strategy that achieves a complex task can be *verified* against an LTL_f spec, or even *synthesized* directly from the spec (which is a form of interpretable policy construction). There is work on extracting logical formulas from learned strategies, or constraining policy search to those that satisfy temporal logic constraints (cf. **shielded RL** and **logic-constrained games**). Overall, the marriage of formal methods and game theory is quite active: it brings the rigor of verification to strategic decision-making, ensuring that automatically learned or synthesized strategies not only *perform* well but also *meet high-level guarantees* (safety, liveness, fairness) expressed in temporal logic ⁸⁰ ⁸¹.

Interpretability in Games (Sparse Strategies & Disentangling Style from Skill)

- **Interpretable Policy Distillation:** As AI agents master complex games (e.g. Go, StarCraft), understanding their strategies becomes crucial. One approach is **policy distillation**: compressing a high-performing (possibly neural network) policy into a simpler, more interpretable form. For example, *Bastani et al. (2018)* extracted decision-tree policies from deep RL agents in Atari games, and *Topin & Veloso (2019)* distilled **explainable symbolic policies**. A recent framework by *Masure et al. (2025)* uses **advantage-based imitation** to distill a deep policy into a linear model ⁸². By incorporating the teacher policy's advantage function into the distillation loss, they ensure the distilled policy focuses on state-action decisions that most affect returns ⁸³. This yields a transparent policy (e.g. a weighted rule or linear combination of features) that remains faithful to the expert's decision-making on critical choices ⁸⁴ ⁸⁵. Empirically, such distilled policies can achieve over 90% of the original performance in games or even financial trading scenarios, while being human-comprehensible ⁸⁶ ⁸⁷. This line of research lies at the intersection of explainable AI and game strategy, aiming to answer “**Tell me about your policy**” in human terms.
- **Sparse and Explainable Strategies:** In multi-agent games, interpretability often means identifying a **small set of key strategies** or **dominant action patterns**. Techniques akin to L1-regularization or feature selection have been applied to strategy profiles. For instance, *Wang et al. (2019)* trained agents in MOBA games and then used **SHAP values** and feature pruning to identify which game states and actions drive agent decisions. In self-play training, one can encourage agents to use a smaller set of distinct tactics (to ease analysis) by adding a complexity penalty. This connects to the earlier discussion of sparse equilibria – beyond computational benefits, sparse strategies are simply easier for humans to interpret (e.g., “this strategy mostly revolves around two opening moves and two responses”). Some research on **alpha-beta pruning of strategy space** and **hierarchical game abstractions** (in poker and war games) effectively yields a more interpretable strategy by eliminating dominated or low-value actions.
- **Style vs. Skill Disentanglement:** In games like chess or StarCraft, players (or agents) can vary in **skill** (effectiveness) and **style** (the manner or pattern of play). Recent work seeks to disentangle these factors using machine learning. One approach is to represent a player's policy as a mixture of a *skill vector* and a *style vector*. *Omnipotent et al. (2021)* used variational autoencoders on game replays to separate a latent dimension correlated with win-rate (skill) from orthogonal dimensions capturing stylistic preferences (e.g. aggressive vs defensive play). Another line, exemplified by *Liu et al. (2023)*, explicitly maintains a **population of self-play agents** with diverse styles but all at high skill levels ²⁴ ⁸⁸. They frame training as a bi-objective optimization: maximize win-rate while also maximizing behavioral diversity among agents. By ensuring only high-performing agents are kept, but among those encouraging diversity, they produced a set of top-tier strategies that are stylistically distinct (e.g. rush tactics, economic focus, stealthy play) ⁸⁹. Such diversity was measured by gameplay feature statistics, and the result was improved generalization — the league of agents could better handle novel opponents because it had

internalized many playstyles. This kind of **evolutionary diverse self-play** was partly inspired by *DeepMind's AlphaStar* league training, which also spawned agents of different “personalities” while keeping them at similar Elo ratings.

- **Behavioral Profiling and Explainable Opponents:** In competitive games, understanding an opponent’s style can be as important as raw skill. Researchers have applied clustering and dimensionality reduction to large datasets of human games (e.g., millions of Go moves or Dota 2 replays) to find archetypal styles. These profiles (like “hyper-aggressive early rush” vs “long-term turtling”) can then be used to explain **why** an agent won or lost (“it struggled against the aggressive style”). They can also guide training: an agent might use style clusters to diversify its training opponents or to adapt strategies mid-game by detecting the opponent’s style cluster.
- **Key Contributions:** A notable paper in this area is *Lan et al. (2020)* “Transparency in Multiagent Games” which introduced **policy graphs** to summarize how agents transition between strategic modes during a game, offering a high-level interpretive graph of strategy flow. Another is *Yeh et al. (2021)* on **game replays to rule induction**, which mined logical if-then rules from StarCraft II replays that roughly describe strategies (e.g. “if scouting sees X, do Y”). While these rules are not perfect, they align with human coaching language.
- **Impact:** Interpretability research in games not only helps humans trust and learn from AI (e.g. a distilled strategy for Go might reveal a novel joseki sequence that humans can adopt), but also aids in debugging and improving AI. If an agent’s strategy is distilled to a few rules, designers can spot irrational or suboptimal rules and refine the training. In competitive settings, an explainable model of an opponent can be used for **transparent model-driven training**, where you deliberately train against a simplified model of the opponent’s strategy to find exploitable weaknesses. As AI systems play larger roles (in economic simulations, autonomous vehicles interacting, etc.), having interpretable game-theoretic strategies will be crucial for **accountability** and **collaboration** between humans and AI. Each of the references above represents a step toward the grand goal of *opening the black box* of game-playing AI and strategic decision models, making their behavior understandable in human terms.

1 2 4 5 8 9 10 12 14 Games with Unawareness

<https://ideas.repec.org/p/ecl/stabus/2122.html>

3 6 7 11 13 15 Unawareness—A gentle introduction to both the literature and the special issue - ScienceDirect

<https://www.sciencedirect.com/science/article/abs/pii/S0165489614000225>

16 20 21 [2412.14337] Commitment to Sparse Strategies in Two-Player Games

<https://arxiv.org/abs/2412.14337>

17 Playing large games using simple strategies - ACM Digital Library

<https://dl.acm.org/doi/10.1145/779928.779933>

18 19 22 23 Computation of Sparse and Dense Equilibrium Strategies of Evolutionary Games

<https://www.mdpi.com/2073-4336/9/3/46>

24 88 89 Explicitly Maintaining Diverse Playing Styles in Self-Play | OpenReview

<https://openreview.net/forum?id=Gkbxt7ThQxU>

25 26 27 cs.toronto.edu

https://www.cs.toronto.edu/~cebly/Papers/sas_mdps_ijcai18.pdf

28 29 30 [PDF] Sparsity-Agnostic Linear Bandits with Adaptive Adversaries - arXiv

<https://arxiv.org/pdf/2406.1192>

31 **Game theory - Wikipedia**

https://en.wikipedia.org/wiki/Game_theory

32 **Game theory - Wikipedia | Encyclopedia - NiNa.Az**

https://www.dl1.en-us.nina.az/Game_theory.html

33 34 35 36 42 43 **web.eecs.umich.edu**

<https://web.eecs.umich.edu/~baveja/Papers/ijcai05.pdf>

37 38 39 40 41 **Optimal Rates for Feasible Payoff Set Estimation in Games**

<https://arxiv.org/html/2602.04397v1>

44 47 48 49 50 57 **Dynamic Scheduling via Polymatroid Optimization | Springer Nature Link**

https://link.springer.com/chapter/10.1007/3-540-45798-4_5

45 46 **ziv.codes**

<https://ziv.codes/pdf/tutorials2025-scully.pdf>

51 52 53 **microsoft.com**

https://www.microsoft.com/en-us/research/wp-content/uploads/2016/04/jmlr16_cmab_fixed.pdf

54 55 56 **Matroid Bandits: Fast Combinatorial Optimization with Learning**

<https://auai.org/uai2014/proceedings/individuals/161.pdf>

58 59 68 69 70 71 72 73 80 81 **LTLf Synthesis with Fairness and Stability Assumptions**

<https://cdn.aaai.org/ojs/5704/5704-13-8929-1-10-20200513.pdf>

60 61 62 63 64 65 74 79 **LTLf: Linear Temporal Logic over Finite Traces**

<https://www.emergentmind.com/topics/linear-temporal-logic-over-finite-traces-ltlf>

66 67 75 76 77 78 **Alternating-time Temporal Logic on Finite Traces**

<https://www.ijcai.org/proceedings/2018/0011.pdf>

82 83 84 85 86 87 **"So, Tell Me About Your Policy...": Distillation of interpretable policies from Deep Reinforcement Learning agents**

<https://arxiv.org/html/2507.07848v1>